

Crop Recommendation

```
import pandas as pd

import matplotlib.pyplot as plt

import seaborn as sns

from sklearn.model_selection import train_test_split

from sklearn.preprocessing import StandardScaler

from sklearn.linear_model import LogisticRegression

from sklearn.neighbors import KNeighborsClassifier

from sklearn.tree import DecisionTreeClassifier, plot_tree

from sklearn.naive_bayes import GaussianNB

from sklearn.svm import SVC

from sklearn.ensemble import RandomForestClassifier, VotingClassifier

from sklearn.metrics import classification_report, accuracy_score, confusion_matrix
```

Ensemble Model

```
# 1. Load the dataset

data = pd.read_csv('Crop_Recommendation.csv') # Adjust path if needed

print(data.head())

# 2. Features and Target

X = data.drop('Crop', axis=1)

y = data['Crop']

# 3. Train-Test Split

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# 4. Feature Scaling (Not mandatory for tree-based models, but okay if used)

scaler = StandardScaler()

X_train_scaled = scaler.fit_transform(X_train)

X_test_scaled = scaler.transform(X_test)

# 5. Train Random Forest Model

rf_model = RandomForestClassifier(

    n_estimators=200,    # Number of trees

    max_depth=None,     # Let it grow fully (or you can set a depth like 10, 20 to control overfitting)

    random_state=42

)

rf_model.fit(X_train, y_train) # No scaling needed for Random Forest (so using X_train directly)
```

```

# 6. Predictions
y_pred = rf_model.predict(X_test)

# 7. Evaluation
print("Accuracy:", accuracy_score(y_test, y_pred))

print("\nClassification Report:\n", classification_report(y_test, y_pred))

# 8. Confusion Matrix
plt.figure(figsize=(12, 8))

sns.heatmap(

    confusion_matrix(y_test, y_pred),

    annot=True, fmt="d", cmap="YlGnBu",

    xticklabels=rf_model.classes_,

    yticklabels=rf_model.classes_

)

plt.title("Confusion Matrix (Random Forest)")

plt.xlabel("Predicted")

plt.ylabel("Actual")

plt.xticks(rotation=45)

plt.yticks(rotation=45)

plt.tight_layout()

plt.show()

```

output

Nitrogen Phosphorus Potassium Temperature Humidity pH_Value \

| | | | | | | |
|---|----|----|----|-----------|-----------|----------|
| 0 | 90 | 42 | 43 | 20.879744 | 82.002744 | 6.502985 |
| 1 | 85 | 58 | 41 | 21.770462 | 80.319644 | 7.038096 |
| 2 | 60 | 55 | 44 | 23.004459 | 82.320763 | 7.840207 |
| 3 | 74 | 35 | 40 | 26.491096 | 80.158363 | 6.980401 |
| 4 | 78 | 42 | 42 | 20.130175 | 81.604873 | 7.628473 |


Rainfall Crop

| | | |
|---|------------|------|
| 0 | 202.935536 | Rice |
| 1 | 226.655537 | Rice |
| 2 | 263.964248 | Rice |
| 3 | 242.864034 | Rice |
| 4 | 262.717340 | Rice |

Accuracy: 0.9931818181818182

Classification Report:

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| Apple | 1.00 | 1.00 | 1.00 | 23 |
| Banana | 1.00 | 1.00 | 1.00 | 21 |
| Blackgram | 1.00 | 1.00 | 1.00 | 20 |
| ChickPea | 1.00 | 1.00 | 1.00 | 26 |
| Coconut | 1.00 | 1.00 | 1.00 | 27 |
| Coffee | 1.00 | 1.00 | 1.00 | 17 |
| Cotton | 1.00 | 1.00 | 1.00 | 17 |
| Grapes | 1.00 | 1.00 | 1.00 | 14 |
| Jute | 0.92 | 1.00 | 0.96 | 23 |
| KidneyBeans | 1.00 | 1.00 | 1.00 | 20 |
| Lentil | 0.92 | 1.00 | 0.96 | 11 |
| Maize | 1.00 | 1.00 | 1.00 | 21 |
| Mango | 1.00 | 1.00 | 1.00 | 19 |
| MothBeans | 1.00 | 0.96 | 0.98 | 24 |
| MungBean | 1.00 | 1.00 | 1.00 | 19 |
| Muskmelon | 1.00 | 1.00 | 1.00 | 17 |
| Orange | 1.00 | 1.00 | 1.00 | 14 |
| Papaya | 1.00 | 1.00 | 1.00 | 23 |
| PigeonPeas | 1.00 | 1.00 | 1.00 | 23 |
| Pomegranate | 1.00 | 1.00 | 1.00 | 23 |
| Rice | 1.00 | 0.89 | 0.94 | 19 |
| Watermelon | 1.00 | 1.00 | 1.00 | 19 |
| accuracy | | 0.99 | | 440 |
| macro avg | 0.99 | 0.99 | 0.99 | 440 |
| weighted avg | 0.99 | 0.99 | 0.99 | 440 |



Support vector Machine

```
X_train,X_test,y_train,y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
# 6. Feature scaling (important for SVM)
```

```
scaler = StandardScaler()
```

```
X_train_scaled = scaler.fit_transform(X_train)
```

```
X_test_scaled = scaler.transform(X_test)
```

```
# 7. Train the SVM model (you can adjust kernel, C, and other parameters)
```

```
model = SVC(kernel='linear') # You can change the kernel to 'rbf', 'poly', etc.
```

```
model.fit(X_train_scaled, y_train)
```

```
# 8. Predictions
```

```
y_pred = model.predict(X_test_scaled)
```

```
# 9. Evaluation
```

```
print("Accuracy:", accuracy_score(y_test, y_pred))
```

```
print("\nClassification Report:\n", classification_report(y_test, y_pred))
```

```
# 10. Confusion Matrix
```

```
plt.figure(figsize=(12, 8))
```

```

sns.heatmap(confusion_matrix(y_test, y_pred), annot=True, fmt="d", cmap="Blues", xticklabels=model.classes_,
yticklabels=model.classes_)

plt.title("Confusion Matrix")

plt.xlabel("Predicted")

plt.ylabel("Actual")

plt.xticks(rotation=45)

plt.yticks(rotation=45)

plt.tight_layout()

plt.show()

```


Output:

Accuracy: 0.9772727272727273

Classification Report:

| | precision | recall | f1-score | support |
|-------------|-----------|--------|----------|---------|
| Apple | 1.00 | 1.00 | 1.00 | 23 |
| Banana | 1.00 | 1.00 | 1.00 | 21 |
| Blackgram | 0.95 | 1.00 | 0.98 | 20 |
| ChickPea | 1.00 | 1.00 | 1.00 | 26 |
| Coconut | 1.00 | 1.00 | 1.00 | 27 |
| Coffee | 1.00 | 1.00 | 1.00 | 17 |
| Cotton | 0.94 | 1.00 | 0.97 | 17 |
| Grapes | 1.00 | 1.00 | 1.00 | 14 |
| Jute | 0.85 | 0.96 | 0.90 | 23 |
| KidneyBeans | 0.91 | 1.00 | 0.95 | 20 |
| Lentil | 0.92 | 1.00 | 0.96 | 11 |
| Maize | 1.00 | 0.95 | 0.98 | 21 |
| Mango | 1.00 | 1.00 | 1.00 | 19 |
| MothBeans | 1.00 | 0.96 | 0.98 | 24 |
| MungBean | 1.00 | 1.00 | 1.00 | 19 |
| Muskmelon | 1.00 | 1.00 | 1.00 | 17 |
| Orange | 1.00 | 1.00 | 1.00 | 14 |
| Papaya | 1.00 | 1.00 | 1.00 | 23 |
| PigeonPeas | 1.00 | 0.87 | 0.93 | 23 |
| Pomegranate | 1.00 | 1.00 | 1.00 | 23 |
| Rice | 0.94 | 0.79 | 0.86 | 19 |
| Watermelon | 1.00 | 1.00 | 1.00 | 19 |
| accuracy | | | 0.98 | 440 |
| macro avg | 0.98 | 0.98 | 0.98 | 440 |

weighted avg 0.98 0.98 0.98 440



Logistic Regression:

3. Train-Test Split

```
X_train,X_test,y_train,y_test = train_test_split(X,y,test_size=0.2,random_state=42)
```

4. Feature Scaling

```
scaler = StandardScaler()
```

```
X_train_scaled = scaler.fit_transform(X_train)
```

```
X_test_scaled = scaler.transform(X_test)
```

5. Train Logistic Regression Model

```
model = LogisticRegression()
```

```
max_iter=300,      # Increase max_iter if convergence warning occurs
```

```
solver='lbfgs',    # Good solver for multinomial problems
```

```
multi_class='multinomial' # Because we have multiple crop classes
```

```
)
```

```
model.fit(X_train_scaled,y_train)
```

6. Predictions

```
y_pred = model.predict(X_test_scaled)
```

7. Evaluation

```
print("Accuracy:", accuracy_score(y_test,y_pred))
```

```
print("\nClassification Report:\n", classification_report(y_test, y_pred))
```

8. Confusion Matrix

```
plt.figure(figsize=(12, 8))
```

```
sns.heatmap(
```

```
confusion_matrix(y_test, y_pred),
```

```
annot=True, fmt="d", cmap="Blues",
```

```
xticklabels=model.classes_,
```

```
yticklabels=model.classes_
```

)

```
plt.title("Confusion Matrix")
```

```
plt.xlabel("Predicted")
```

```
plt.ylabel("Actual")
```

```
plt.xticks(rotation=45)
```


```
plt.yticks(rotation=45)
```

```
plt.tight_layout()
```

```
plt.show()
```

Output:

Accuracy: 0.9636363636363636



KNN Model:

```
# 3. Train-Test Split

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# 4. Feature Scaling (very important for KNN)

scaler = StandardScaler()

X_train_scaled = scaler.fit_transform(X_train)

X_test_scaled = scaler.transform(X_test)

# 5. Train KNN Model

knn_model = KNeighborsClassifier(n_neighbors=5) # You can tune n_neighbors later

knn_model.fit(X_train_scaled, y_train)

# 6. Predictions

y_pred = knn_model.predict(X_test_scaled)

# 7. Evaluation

print("Accuracy:", accuracy_score(y_test, y_pred))

print("\nClassification Report:\n", classification_report(y_test, y_pred))

# 8. Confusion Matrix

plt.figure(figsize=(12, 8))

sns.heatmap(

    confusion_matrix(y_test, y_pred),

    annot=True, fmt="d", cmap="Greens",

    xticklabels=knn_model.classes_,

    yticklabels=knn_model.classes_

)

plt.title("Confusion Matrix (KNN)")

plt.xlabel("Predicted")

plt.ylabel("Actual")

plt.xticks(rotation=45)

plt.yticks(rotation=45)

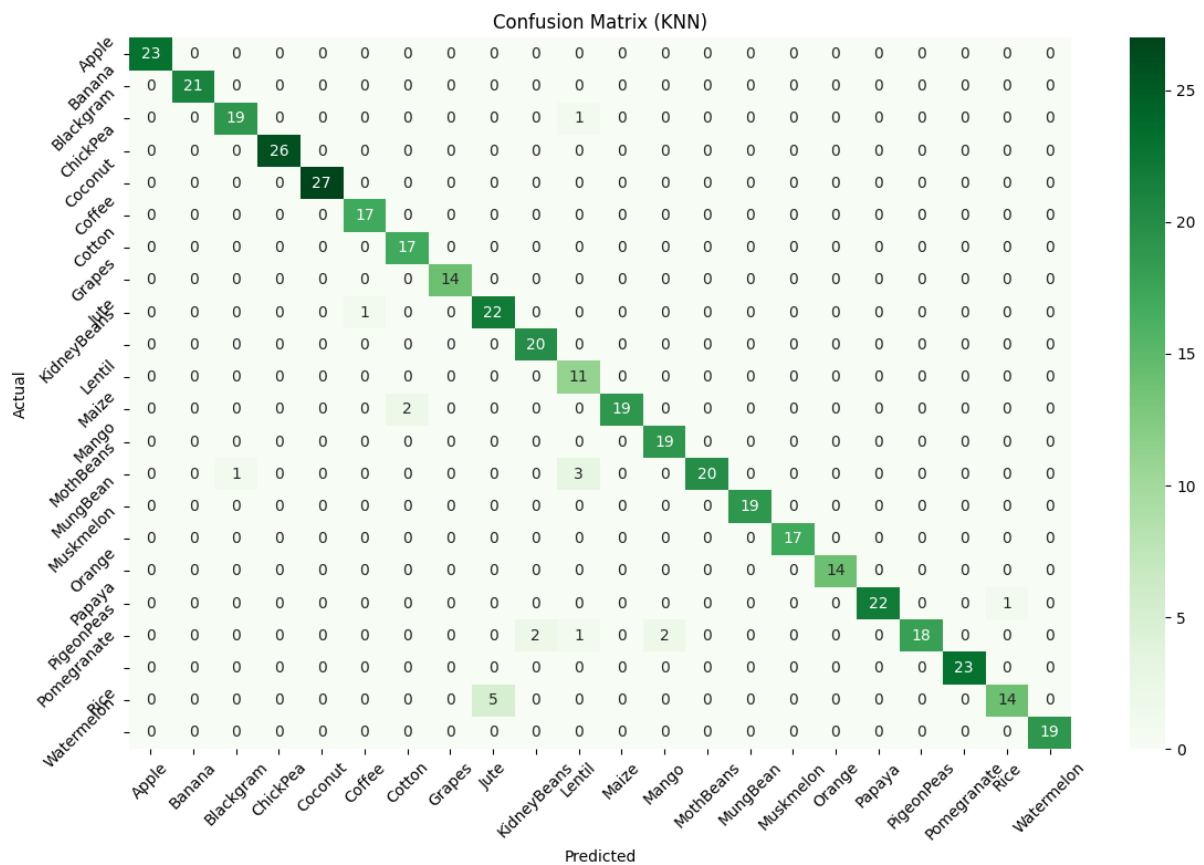
plt.tight_layout()

plt.show()
```

Output:

Accuracy: 0.9568181818181818

| | | | | |
|--------------|------|------|------|-----|
| accuracy | | | 0.96 | 440 |
| macro avg | 0.96 | 0.96 | 0.95 | 440 |
| weighted avg | 0.96 | 0.96 | 0.96 | 440 |



Decision Tree:

2. Features and Target

```
X = data.drop('Crop', axis=1)
```

```
y = data['Crop']
```

3. Train-Test Split

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

4. Feature Scaling (NOT necessary for Decision Trees, but safe if you plan to combine later)

```
scaler = StandardScaler()
```

```
X_train_scaled = scaler.fit_transform(X_train)
```

```
X_test_scaled = scaler.transform(X_test)
```

5. Train Decision Tree Model

```
dt_model = DecisionTreeClassifier(random_state=42, max_depth=None) # You can tune max_depth if needed
```

```
dt_model.fit(X_train, y_train) # For tree models, scaling not necessary. So use X_train directly.
```

6. Predictions

```
y_pred = dt_model.predict(X_test)
```

7. Evaluation

```
print("Accuracy:", accuracy_score(y_test, y_pred))
```

```
print("\nClassification Report:\n", classification_report(y_test, y_pred))
```

8. Confusion Matrix

```
plt.figure(figsize=(12, 8))

sns.heatmap(

    confusion_matrix(y_test, y_pred),

    annot=True, fmt="d", cmap="Oranges",

    xticklabels=dt_model.classes_,

    yticklabels=dt_model.classes_

)

plt.title("Confusion Matrix (Decision Tree)")

plt.xlabel("Predicted")

plt.ylabel("Actual")

plt.xticks(rotation=45)

plt.yticks(rotation=45)

plt.tight_layout()

plt.show()

plt.figure(figsize=(20, 10))

plot_tree(dt_model, filled=True, feature_names=X.columns, class_names=dt_model.classes_, fontsize=10)

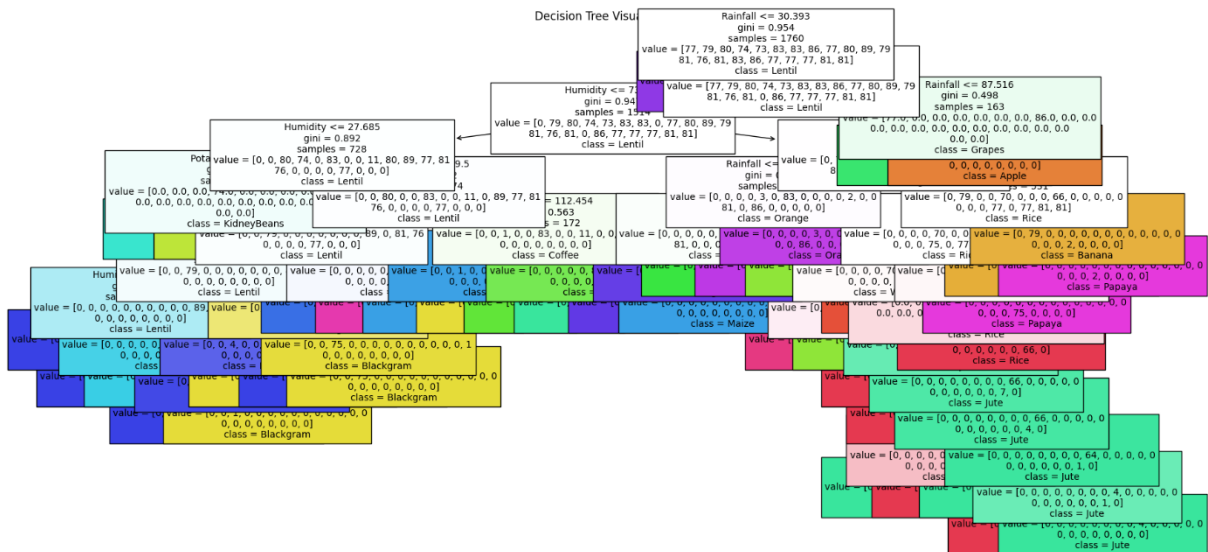

plt.title("Decision Tree Visualization")

plt.show()
```

Output:

Accuracy: 0.9863636363636363

| | | | | | |
|--------------|------|------|------|------|-----|
| accuracy | | | | 0.99 | 440 |
| macro avg | 0.99 | 0.99 | 0.99 | 0.99 | 440 |
| weighted avg | 0.99 | 0.99 | 0.99 | 0.99 | 440 |



Naïve Base:

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

4. Feature Scaling

```
scaler = StandardScaler()
```

```
X_train_scaled = scaler.fit_transform(X_train)
```

```
X_test_scaled = scaler.transform(X_test)
```

5. Train Naive Bayes Model

```
nb_model = GaussianNB()
```

```
nb_model.fit(X_train_scaled, y_train)
```

6. Predictions

```
y_pred = nb_model.predict(X_test_scaled)
```

7. Evaluation

```
print("Accuracy:", accuracy_score(y_test, y_pred))
```

```
print("\nClassification Report:\n", classification_report(y_test, y_pred))
```

8. Confusion Matrix

```
plt.figure(figsize=(12, 8))
```

```
sns.heatmap(
```

```
    confusion_matrix(y_test, y_pred),
```

```
    annot=True, fmt="d", cmap="Purples",
```

```
    xticklabels=nb_model.classes_,
```

```
    yticklabels=nb_model.classes_
```

```
)
```

```
plt.title("Confusion Matrix (Naive Bayes)")
```

```
plt.xlabel("Predicted")
```

```
plt.ylabel("Actual")
```

```
plt.xticks(rotation=45)
```

```
plt.yticks(rotation=45)
```

```
plt.tight_layout()
```

```
plt.show()
```

Output:

```
Accuracy: 0.9954545454545455
```

| | | | | | | |
|--------------|------|------|------|-----|------|-----|
| accuracy | | | | | 1.00 | 440 |
| macro avg | 1.00 | 1.00 | 1.00 | 440 | | |
| weighted avg | 1.00 | 1.00 | 1.00 | 440 | | |

